

African Music Network Dataset Project

Mark Lekina Rorat

June 5, 2023

https://github.com/marklekina/african_music_dataset_project/

1 Introduction

This report presents a new network dataset of African music artists and track data. We demonstrate how we collected artist and track metadata using the Spotify API, and artist nationality data through web scraping of the African musicians Wikipedia page. We also conducted some investigations on the dataset, namely the correlation between the proportion of collaborations and artist nationality/genre, the evolution of collaboration patterns over time, and the relationship between collaboration and artist popularity. Through statistical analyses, we provide insights into the African music industry and potential opportunities for fostering collaborations and promoting African artists on a global scale. First, we describe our dataset through the data collection and preprocessing workflow. Next, we present the results of our analyses coupled with a summary of our findings. Finally, we conclude with the limitations of our work and suggestions for future research.

2 Dataset overview

2.1 Data sources

To build our network dataset, we utilized these two data sources:

- *Spotify* API¹: This API provides extensive artist and track data spanning various genres and time periods. We used the API to gather valuable information such as track popularity, artist follower counts, and other relevant metadata.
- *Wikipedia*²: Despite its breadth, the Spotify API does not include artist nationality data. For this reason, we scraped the Wikipedia page listing African musicians to collect nationality information for

¹<https://developer.spotify.com/>

²https://en.wikipedia.org/wiki/List_of_African_musicians

the artists.

2.2 Data collection

During the data collection phase, we utilized two scripts to gather the necessary information:

- **fetch_spotify_data:** This script was responsible for querying the Spotify API to retrieve artist and track data. Using the `spotifyr` library, we made API calls and obtained JSON responses. The script then parsed these responses using the `jsonlite` library, converting them into structured data frames for further analysis.
- **scrape_nationality_data:** This script focused on crawling artist Wikipedia pages to extract information about their nationalities. Using the `rvest` library, we performed web scraping to collect the desired data. The script compiled the nationality information into a central dataset, which would later be merged with the Spotify data.




Figure 1: Network of African Artists by Country

2.3 Data collection workflow

The data collection workflow involved several steps to gather the required information for our analysis. Here is an overview of the process:

1. Starting with the base URL for African musicians on Wikipedia, we fetched the HTML content from webpages containing lists of artist names grouped by country.
2. We performed web scraping on the base webpage to extract the artist names and country information. These details were stored in a data frame.
3. For countries that had links to webpages with artist names instead of directly listing them, we navigated to those pages and scraped the artist names into the data frame.
4. Using the artist names collected from Wikipedia, we utilized the Spotify API's search endpoint to fetch artist objects for each name. We only considered exact matches to ensure relevance.
5. To supplement our list of artists, we retrieved the related artists for each artist using the Spotify API's related artists endpoint.
6. After gathering all the necessary artist objects, we parsed the artist metadata from the JSON responses and stored it in a data frame.
7. To collect track data, which was required for aggregating collaboration information, we employed two approaches:
 - First, we obtained the top tracks for each artist using the appropriate API endpoint.
 - Second, we searched for track objects by artist name. We then subsetted the results to include only tracks where the artist ID matched the ID provided in the track's list of artists.
8. Once we had collected all the track JSON objects, we parsed the track metadata into a separate data frame.

2.4 Data preprocessing workflow

After data collection, we did some data preprocessing to prepare the collected data for analysis. For this part, we used the `preprocessing_analysis` script to perform the data preprocessing tasks and some preliminary analysis. Here is an overview of the process:

1. The first step was merging the artist Spotify data with the nationality data collected from Wikipedia.

2. Next, we addressed the issue of duplicate tracks in the track data. Since each track has a unique ID, we couldn't deduplicate based on ID alone. Instead, we compared the list of artists performing the track and the raw version of the track's name after removing whitespace and matching cases. If both matched, we selected the track with the highest popularity and dropped the duplicates. Typically, the most popular track corresponds to the original album release, while the duplicates may be compilation releases or versions not by the artist.
3. The next step involved building the network adjacency matrix from the track and artist data. We first subsetting the data to include only tracks with more than one artist. Each artist in these tracks is considered a *node* in the network, and an *edge* represents a collaboration between two artists. The *weight* of each edge is the number of collaborations between the artists in the track dataset.
4. We pruned the network to include only edges between known African artists. This is because some of our analysis requires nationality data, which is not available for each node. From this subset, we induced a subgraph from the largest connected component of the graph. The final result was a graph with 342 *nodes* and 1111 *edges*, spanning 23 *countries*. Figure 1 visualizes the network. Each node is colored to match the artist's nationality, and the colors for the countries with the highest number of nodes in the network are indicated in the legend.

2.5 Strengths of the dataset

One of the strengths of this dataset is that it represents the first comprehensive effort to aggregate African music data into a coherent and publicly available dataset. It brings together information from multiple sources, including the Spotify API and Wikipedia, to provide a comprehensive view of African music artists and their music. The dataset incorporates nationality data from Wikipedia, which enriches the artist and track metadata obtained from Spotify. This additional information allows for a more comprehensive understanding of the artists and their cultural backgrounds.

Another strength of the dataset is its balanced representation of artists from different African countries. Efforts were made to ensure that the dataset is not biased towards artists from countries with more popular music scenes. The Spotify API was queried for data on artists from each African country listed on the Wikipedia page, ensuring a diverse representation of artists. Furthermore, the transparency of the dataset is a notable strength. The code used for data collection, preprocessing, and analysis is publicly available, allowing for reproducibility and modification. This openness provides opportunities for researchers to improve the dataset and enhance its quality over time.

2.6 Weaknesses of the dataset

One of the main weaknesses of the dataset is the challenge of obtaining comprehensive and accurate data on artist nationalities. The availability of extensive and high-quality nationality data for artists is crucial for conducting meaningful analyses. However, this information is often incomplete or inaccurate. One challenge is that the Wikipedia web pages on African artists cannot possibly list all artists from each specific country. As a result, it is impossible to fetch artist metadata for a significant majority of artists from a particular country. This limitation restricts the dataset’s coverage and may result in an incomplete representation of artists from certain countries.

Additionally, some Wikipedia pages may incorrectly list the nationality of certain artists. This is particularly common for artists who have gained global popularity but are not necessarily of African nationality, but have African ancestral roots. Detecting and correcting these inaccuracies at scale is a difficult task. As a result, relying on nationality data to diversify the dataset and conduct further analysis becomes challenging, as low-quality nationality data can undermine the accuracy of the findings.

To address this issue, the dataset had to be subsetting to include only artists whose nationalities are Wikipedia-verified. This necessary step in ensuring data quality and accuracy results in a significant loss of data that could have otherwise been valuable for analysis. Consequently, the final network dataset includes only 23 nationalities, despite starting data preprocessing with artists from over 45 different nationalities. Overall, the limitations in obtaining comprehensive and accurate artist nationality data pose a significant weakness in the dataset. It highlights the challenges faced in compiling a comprehensive dataset for African music and the need for further efforts to improve data quality and coverage.

3 Analysis and Findings

The key questions we wanted to investigate for our analysis are the following:

1. How does proportion of collaborations vary by country?
2. Which genres tend to have the highest (or the lowest) proportion of collaborations?
3. How has collaboration between African artists evolved over time?
4. Is there a correlation between collaboration and the popularity of artists?

3.1 Preliminary analysis

We addressed the first three questions through straightforward analysis methods:

3.1.1 Collaboration Proportion by Artist Nationality

country	collaboration_proportion	collaboration_count
Swaziland	0.7435897	29
South Sudan	0.5671642	38
Tanzania	0.4983165	148
Benin	0.4285714	30
Nigeria	0.3707527	1724
Ghana	0.3647235	732
Kenya	0.3397933	263
Democratic Republic of the Congo	0.3146269	527
Cape Verde	0.2888889	13
Mali	0.2820513	165

Table 1: Top 10 countries with the highest proportion of collaborative tracks

To analyze the proportion of collaborations based on artist nationality, we merged the track data with artist nationality data and calculated the number of rows (*i.e.*, tracks) where each unique country name appeared for both singles and collaborations. The collaboration proportion was computed as the ratio of collaboration count to the total count. The countries with the highest and lowest collaboration proportions are presented in Tables 1 and 2. As expected, countries with a larger number of artists in the network tended to have higher collaboration proportions.

country	single_artist_proportion	single_artist_count
Eritrea	1.0000000	46
Guinea-Bissau	0.9772727	43
Togo	0.9759036	81
Sudan	0.9743590	38
Sierra Leone	0.9655172	56
Gabon	0.9647059	82
Botswana	0.9245283	49
Egypt	0.9184028	529
Guinea	0.8983051	159
Republic of the Congo	0.8673469	170

Table 2: Top 10 countries with the highest proportion of single-artist tracks

3.1.2 Collaboration Proportion by Artist Genres

genre	collaboration_proportion	collaboration_count
amapiano	1.0000000	13
baroque singing	1.0000000	10
british modern classical	1.0000000	28
early modern classical	1.0000000	28
impressionism	1.0000000	28
post-romantic era	1.0000000	28
sgija	1.0000000	13
gqom	0.9729730	36
igbo traditional	0.9354839	29
edm	0.8604651	37

Table 3: Top 10 genres by collaboration proportion

genre	single_artist_proportion	single_artist_count
african gospel	1	42
black punk	1	34
botswana traditional	1	43
congolese gospel	1	42
gospel amapiano	1	3
microtonal	1	39
ngoni	1	43
oromo pop	1	9
sierra leonean pop	1	14
south african metal	1	17

Table 4: Top 10 genres by single-artist proportion

For this analysis, we utilized genre data available in the artist Spotify metadata. The metadata includes a list of genres commonly associated with an artist’s music. Similar to the nationality analysis, we computed the collaboration proportions for each genre by merging track data with artist genre data. The results are shown in Tables 3 and 4. Our qualitative analysis revealed that non-native genres, such as classical and impressionism music, tend to feature more collaborations, while country-specific indigenous music often leans towards singles.

3.1.3 Collaboration Over Time




Figure 2: Proportion of collaborative tracks over time

To investigate the temporal evolution of collaboration, we grouped the track data by the release year of albums (extracted from the album’s release date in track metadata) and computed the collaboration proportions. Figure 2 presents a plot of collaboration proportions of the results over time. Although the proportion of collaborations exhibits a peak in the mid-70s, the overall volume of collaborations tends to increase over time. However, it is important to note that this observation may be influenced by the larger presence of recent tracks in the dataset, and further analysis is needed to draw definitive conclusions.

By conducting these analyses, we aimed to gain insights into the correlation between artist popularity and artist network centrality measures. These findings provide a foundation for understanding the dynamics of collaboration within the African music industry.

3.2 Network analysis

3.2.1 What is the significance of the network's centrality measures in predicting artist popularity?

centrality_measure	term	estimate	std.error	statistic	p.value
betweenness	(Intercept)	3.531621e+01	0.9489058	37.217826	0.0e+00
betweenness	value	4.367100e-03	0.0007625	5.727709	0.0e+00
closeness	(Intercept)	1.821336e+01	3.8605889	4.717766	3.5e-06
closeness	value	3.008932e+04	5786.9529144	5.199510	3.0e-07
degree	(Intercept)	3.156899e+01	1.0988871	28.728144	0.0e+00
degree	value	9.578788e-01	0.1149585	8.332387	0.0e+00
eigenvector	(Intercept)	3.554460e+01	0.9401270	37.808296	0.0e+00
eigenvector	value	4.311637e+01	7.8108501	5.520061	1.0e-07
page_rank	(Intercept)	2.967162e+01	1.3210111	22.461298	0.0e+00
page_rank	value	2.777305e+03	355.6898172	7.808221	0.0e+00

Table 5: Linear Regression of Artist Popularity on Individual Centrality Measures

In order to examine the relevance of the network's centrality measures in predicting artist popularity, we calculated various centrality measures for the network and conducted linear regression analysis. The centrality measures served as independent variables, while artist popularity was the dependent variable. The results of the analysis are presented in Table 5.

The analysis reveals that higher centrality measures tend to correspond to higher artist popularity, as indicated by the consistently higher intercept estimates. The p-values associated with each centrality measure are extremely close to zero, suggesting a high level of statistical significance for each measure.

These findings highlight the importance of centrality measures in predicting artist popularity within the network. Artists who exhibit higher centrality, as indicated by measures such as betweenness, closeness, degree, eigenvector, and page rank, are more likely to have a greater level of popularity. The strong statistical significance of these centrality measures further supports their relevance in understanding and predicting artist popularity within the network.

3.2.2 How well do community structures in the network predict artist popularity?

term	estimate	std.error	statistic	p.value
(Intercept)	45.9821075	1.2118217	37.94461	0
community	-0.7030323	0.0589465	-11.92661	8.6e-33

Table 6: Logistic Regression of Artist Popularity on Community Membership

To assess the predictive power of community structures in the network on artist popularity, we utilized the *Girvan-Newman algorithm* to detect communities based on *edge-betweenness centrality*. The resulting community structures are visually represented in Figure 3.

Since the community membership variable derived from the algorithm is categorical rather than numerical, we applied *one-hot encoding* and employed logistic regression to build a model. In this model, artist popularity served as the dependent variable, while community membership was the independent variable. The outcomes of the logistic regression analysis are presented in Table 3.2.2.

The results demonstrate that community membership significantly influences an artist’s popularity. The p-values, which are close to zero, indicate a highly significant relationship between an artist’s community membership and their level of popularity. Figure 3 provides a qualitative examination of the community structures, revealing that while artist nationality predominantly defines the communities, there are also instances where communities extend beyond national boundaries in the denser parts of the network.

Considering the previous findings that centrality measures, including node betweenness centrality, exhibit a statistically significant relationship with artist popularity, it follows that community memberships defined by edge-betweenness centrality also possess a significant association with an artist’s popularity. This emphasizes the importance of community structures in understanding and predicting artist popularity within the network.




Figure 3: Network of African Artists by Communities

4 Conclusion

In conclusion, this report has presented a comprehensive analysis of African music artists and their collaborations using a dataset that combines information from Spotify and Wikipedia. Through the careful collection and preprocessing of data, we have been able to explore various aspects of collaboration within the African music industry and draw meaningful insights.

Our analysis revealed several interesting findings. Firstly, we observed variations in the proportion of collaborations based on artist nationality and genre. Non-native genres tended to feature more collaborations, while country-specific indigenous music often leaned towards singles. This highlights the diverse nature of African music and the influence of cultural and genre-specific factors on collaboration patterns.

Additionally, we investigated the evolution of collaboration over time and observed a general increasing trend in the volume of collaborations. While collaboration reached its peak in the mid-70s, the presence of recent tracks in the dataset may have skewed this observation. Nevertheless, this analysis provides valuable

context to understand the historical development of collaborations within the African music industry.

Furthermore, our study explored the relationship between collaboration and artist popularity using centrality measures and community structures within the network. The results indicated that higher centrality measures tend to be associated with higher artist popularity, suggesting the significance of network centrality in predicting artist success. Additionally, community membership, identified through the Girvan-Newman algorithm, emerged as a significant predictor of artist popularity, highlighting the influence of community structures on an artist’s recognition and reach.

It is important to acknowledge the limitations of our analysis. One key challenge was the availability and quality of artist nationality data, which required us to subset the dataset to Wikipedia-verified nationalities, resulting in a loss of valuable data. Additionally, the dataset may not capture the entire African music landscape, as it relies on the availability of artist data on Spotify and Wikipedia.

Nevertheless, this study represents an important step in understanding and appreciating the vibrant collaborations within the African music industry. The findings can inform stakeholders in the music industry, policymakers, and researchers interested in African music by providing valuable insights into the dynamics of collaboration and its impact on artist success. Moving forward, we hope to continue expanding and refining the dataset, incorporating data from additional sources and exploring other dimensions of collaboration, such as cross-border collaborations and the influence of social and cultural factors.

Overall, this analysis contributes to the growing body of research on African music and underscores the richness, diversity, and potential of the African music industry. By embracing collaborations and leveraging the unique cultural heritage of African music, we can celebrate its global impact and foster its continued growth and recognition.